

The R logo, consisting of a blue 'R' inside a gray circle.

Artificial intelligence



A guide to using R for ML

Nikhita Damaraju

06-22-2024

Cascadia R conference 2024



Examples of ML



When is ML used?

- High-dimensional data

More features/variables than rows increases complexity of data.





When is ML used?

- High-dimensional data

More features/variables than rows increases complexity of data.

- Non-linear relationships

Indirect relationship between predictors and outcome.



When is ML used?

- High-dimensional data



More features/variables than rows increases complexity of data.

- Non-linear relationships

Indirect relationship between predictors and outcome.



- Unstructured data

Non-tabular data like text, images, audio.



When is ML used?

- **High-dimensional data**
More features/variables than rows increases complexity of data.
- **Non-linear relationships**
Indirect relationship between predictors and outcome.
- **Unstructured data**
Non-tabular data like text, images, audio.
- **Scalability**
Ability to handle millions or billions of data points.



Disclaimer!

This talk will **not** cover:

- R vs python for ML
- How to choose ML methods
- Pros and cons of using R for ML

Building an ML project


Processing data



Picking and training your ML model



Testing the performance of ML models



Building an ML project


Processing data



Picking and training your ML model




Testing the performance of ML models



Data processing

- Exploratory data analysis
- Handle missing data
- Outliers
- Feature engineering
- Feature scaling
- Splitting data into training and validation sets



```
# Example
library(dplyr)
library(ggplot2)

# Summary statistics
summary(mtcars)

# Visualization
ggplot(mtcars, aes(x=mpg, y=hp)) + geom_point()
```

Data processing

- Exploratory data analysis
- Handle missing data
- Outliers
- Feature engineering
- Feature scaling
- Splitting data into training and validation sets



stekhoven/ missForest



missForest is a nonparametric, mixed-type imputation method for basically any type of data for the statistical software R.

3 Contributors 11 Issues 87 Stars 23 Forks

```
# Example
library(mice)
library(missForest)


# Remove rows with missing values
clean_data <- na.omit(mtcars)

# Multiple imputation
imputed_data <- mice(mtcars, m=5, maxit=50, method='pmm', seed=500)

# Imputation using Random Forests
imputed_data_rf <- missForest(mtcars)
```

Data processing

- Exploratory data analysis
- Handle missing data
- Outliers
- Feature engineering
- Feature scaling
- Splitting data into training and validation sets





```
# Example
library(dplyr)
library(ggplot2)

# Identify outliers using boxplot
boxplot(mtcars$mpg)

# Remove outliers
clean_data <- mtcars %>% filter(mpg < quantile(mpg, 0.95))
```

Data processing

- Exploratory data analysis
- Handle missing data
- Outliers
- Feature engineering
- Feature scaling
- Splitting data into training and validation sets



```
# Example
library(dplyr)
library(caret)

# Create new variable
mtcars <- mtcars %>% mutate(power_to_weight = hp / wt)

# Create dummy variables
dummies <- dummyVars(~., data=mtcars)
mtcars_dummies <- predict(dummies, newdata=mtcars)
```

Data processing

- Exploratory data analysis
- Handle missing data
- Outliers
- Feature engineering
- **Feature scaling**
- Splitting data into training and validation sets




```
# Example
library(caret)
library(scales)

# Standardize variables
scaled_data <- scale(mtcars)

# Pre-process using caret
preProc <- preProcess(mtcars, method=c("center", "scale"))
mtcars_scaled <- predict(preProc, mtcars)|
```

Data processing

- Exploratory data analysis
- Handle missing data
- Outliers
- Feature engineering
- Feature scaling
- Splitting data into training and validation sets



cran/caTools

! This is a read-only mirror of the CRAN R package repository. caTools — Tools: Moving Window Statistics, GIF, Base64,...



2 Contributors 0 Issues 7 Stars 3 Forks



```
# Example
library(caret)
library(caTools)

# Using caret
set.seed(123)
trainIndex <- createDataPartition(mtcars$mpg, p=0.7, list=FALSE)
train_data <- mtcars[trainIndex,]
test_data <- mtcars[-trainIndex,]

# Using caTools
set.seed(123)
split <- sample.split(mtcars$mpg, SplitRatio=0.7)
train_data <- subset(mtcars, split==TRUE)
test_data <- subset(mtcars, split==FALSE)
```


Building an ML project


Processing data



Picking and training your ML model



Testing the performance of ML models



Types of ML algorithms

Supervised learning

- algorithm is trained on a labeled dataset
- each input is paired with the correct output

Unsupervised learning



- algorithm works with unlabeled data
- finds patterns or structures
- no explicit guidance

Types of ML algorithms

Supervised learning

Regression



Classification



Unsupervised learning

Clustering

Dimensionality reduction



Supervised learning

Supervised learning

Regression

- **glmnet**: For regularized generalized linear models (lasso, ridge, elastic net)
- **randomForest** or **ranger**: For random forest regression
- **xgboost**: For gradient boosting regression
- **rpart**: For decision tree regression
- **kernlab**: For support vector regression

```
#glmnet: For regularized generalized linear models (lasso, ridge, elastic net)
# Example data
x <- matrix(rnorm(100*20), 100, 20)
y <- rnorm(100)
```

```
# Fit a lasso model
fit <- glmnet(x, y, alpha = 1)
```

```
#randomForest: For random forest regression
# Example data
data(iris)
set.seed(42)
```

```
# Fit a random forest model
fit <- randomForest(Sepal.Length ~ ., data = iris)
```

```
#xgboost: For gradient boosting regression
# Example data
data(iris)
set.seed(42)
train <- as.matrix(iris[, -1])
label <- iris$Sepal.Length
```

```
# Fit a gradient boosting model
fit <- xgboost(data = train, label = label, nrounds = 100, objective = "reg:squareerror")
```

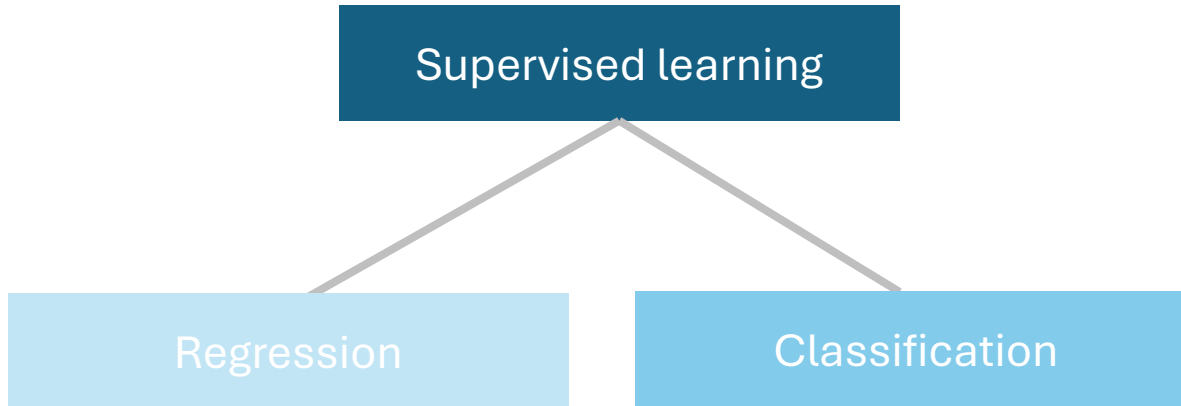
```
#rpart: For decision tree regression
# Example data
data(iris)
```

```
# Fit a decision tree model
fit <- rpart(Sepal.Length ~ ., data = iris)
```

```
#kernlab: For support vector regression
# Example data
data(iris)
```

```
# Fit a support vector regression model
fit <- ksvm(Sepal.Length ~ ., data = iris, type = "eps-svr")
```

Supervised learning



- **caret**: For various classification algorithms and model training
- **randomForest**: For random forest classification
- **xgboost**: For gradient boosting classification
- **e1071**: For support vector machines and naive Bayes
- **rpart**: For decision tree classification
- **glmnet**: For regularized logistic regression
- **naivebayes**: For naive Bayes classification

```
#caret: For various classification algorithms and model training
# Example data
data(iris)
set.seed(42)
```

```
# Train a model using caret
fit <- train(Species ~ ., data = iris, method = "rf")
```

```
#randomForest: For random forest classification
# Example data
data(iris)
set.seed(42)
```

```
# Fit a random forest model
fit <- randomForest(Species ~ ., data = iris)
```

```
#xgboost: For gradient boosting classification
# Example data
data(iris)
set.seed(42)
train <- as.matrix(iris[, -5])
label <- as.numeric(iris$Species) - 1
```

```
# Fit a gradient boosting model
fit <- xgboost(data = train, label = label, nrounds = 100, objective = "multi:softprob")
```

```
#e1071: For support vector machines and naive Bayes
# Example data
data(iris)
```


```
# Fit a support vector machine model
fit_svm <- svm(Species ~ ., data = iris)
```

```
# Fit a naive Bayes model
fit_nb <- naiveBayes(Species ~ ., data = iris)
```

```
#rpart: For decision tree classification
# Example data
data(iris)
```

```
# Fit a decision tree model
fit <- rpart(Species ~ ., data = iris)
```

Unsupervised learning



- **stats:** For k-means clustering (built-in R package)
- **cluster:** For various clustering algorithms including hierarchical clustering
- **dbscan:** For density-based clustering

```
#stats: For k-means clustering (built-in R package)
# Example data
data(iris)

# Perform k-means clustering
fit <- kmeans(iris[, -5], centers = 3)

#cluster: For various clustering algorithms including hierarchical clustering
# Install and load the package
install.packages("cluster")
library(cluster)

# Example data
data(iris)

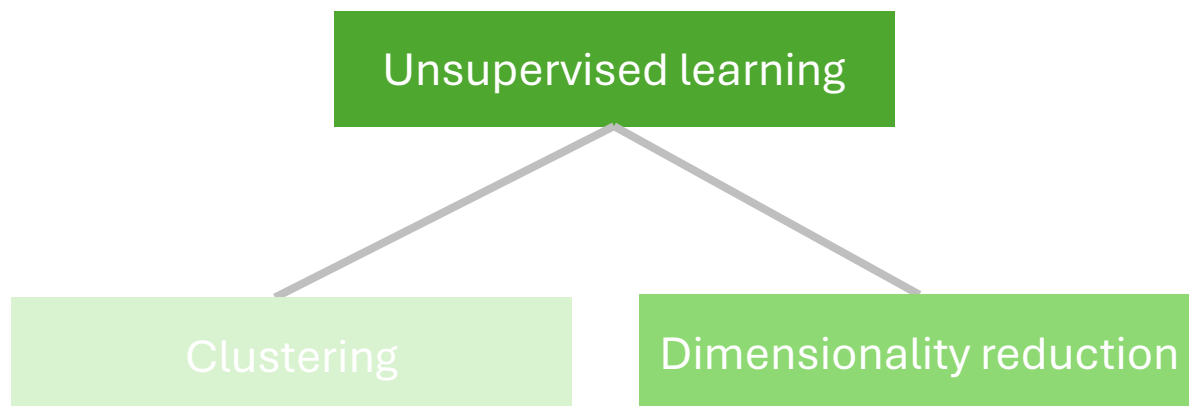
# Perform hierarchical clustering
fit <- agnes(iris[, -5])

#dbscan: For density-based clustering
# Install and load the package
install.packages("dbscan")
library(dbscan)

# Example data
data(iris)

# Perform DBSCAN clustering
fit <- dbscan(iris[, -5], eps = 0.5, minPts = 5)
```

Unsupervised learning



- **stats**: For principal component analysis (PCA) (built-in R package)
- **FactoMineR**: For various dimensionality reduction techniques
- **Rtsne**: For t-SNE (t-Distributed Stochastic Neighbor Embedding)

```
#stats: For principal component analysis (PCA) (built-in R package)
# Example data
data(iris)

# Perform PCA
fit <- prcomp(iris[, -5], scale. = TRUE)

#FactoMineR: For various dimensionality reduction techniques
# Install and load the package
install.packages("FactoMineR")
library(FactoMineR)

# Example data
data(iris)

# Perform PCA using FactoMineR
fit <- PCA(iris[, -5], graph = FALSE)

# Rtsne: For t-SNE (t-Distributed Stochastic Neighbor Embedding)
# Install and load the package
install.packages("Rtsne")
library(Rtsne)

# Example data
data(iris)

# Perform t-SNE
fit <- Rtsne(as.matrix(iris[, -5]))
```


Types of ML algorithms

Supervised learning

Regression

glmnet
randomForest or **ranger**
xgboost
rpart
kernlab

Classification

caret
randomForest
xgboost
e1071
rpart
glmnet
naivebayes

Unsupervised learning

Clustering

stats
cluster
dbscan

Dimensionality reduction

stats or **PCATools**
FactoMineR
Rtsne

Building an ML project


Processing data



Picking and training your ML model



Testing the performance of ML models



Testing ML models

- **Classification Metrics**

- Accuracy using **caret**
- Precision and Recall using **caret**
- F1 Score using **caret**
- AUC-ROC using **pROC** or **ROCit**

- **Regression Metrics**

- RMSE (Root Mean Squared Error)
- R-squared (R^2)

- **Visualization Techniques**

- ROC Curves
- Density Plots and Boxplots

```
# Load necessary libraries
library(caret)
library(pROC)

# Load dataset
data(mtcars)

# Split data into training and testing sets
set.seed(123)
trainIndex <- createDataPartition(mtcars$mpg, p = .8,
                                   list = FALSE,
                                   times = 1)

trainData <- mtcars[ trainIndex,]
testData  <- mtcars[-trainIndex,]

# Make predictions
predictions_prob <- predict(model, testData, type = "response")
predictions <- ifelse(predictions_prob > 0.5, 1, 0)
predictions <- as.factor(predictions)

# Create a confusion matrix
conf_matrix <- confusionMatrix(predictions, testData$Species, positive = "1")
print(conf_matrix)

# Extract metrics
accuracy <- conf_matrix$overall['Accuracy']
precision <- conf_matrix$byClass['Pos Pred Value']
recall <- conf_matrix$byClass['Sensitivity']
f1_score <- 2 * (precision * recall) / (precision + recall)

# ROC Curve
roc_obj <- roc(testData$Species, predictions_prob)
plot(roc_obj)
```

Conclusion


Processing data



Picking and training your
ML model



Testing the performance
of ML models



Thank you